

# assignment4

September 21, 2023

```
[1]: from google.colab import files
      uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving winequality-red.csv to winequality-red.csv

```
[10]: import pandas as pd
      import numpy as np
      import matplotlib.pyplot as plt
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score, classification_report
      from scipy import stats

      # Task 1: Load the Dataset
      data = pd.read_csv("winequality-red.csv")

      # Task 2: Data Preprocessing and Feature Engineering
      # Feature Engineering (Adding a new feature: total acidity)
      data['total_acidity'] = data['fixed acidity'] + data['volatile acidity']

      # Task 3: Data Preprocessing - Removing Outliers using z-score
      z_scores = np.abs(stats.zscore(data))
      threshold = 3
      data_no_outliers = data[(z_scores < threshold).all(axis=1)]

      # Visualize outliers using box plots
      fig, axes = plt.subplots(nrows=4, ncols=3, figsize=(15, 15))
      for i, ax in enumerate(axes.flat):
          if i < len(data.columns):
              ax.boxplot(data_no_outliers[data.columns[i]])
              ax.set_title(data.columns[i])

      plt.tight_layout()
      plt.show()
```

```

# Visualize the histograms after removing outliers
data_no_outliers.hist(bins=20, figsize=(15, 10))
plt.show()

# Split the data into features (X) and target (y)
X = data_no_outliers.drop("quality", axis=1)
y = data_no_outliers["quality"]

# Task 3 cont'd: Data Preprocessing
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Task 4: Machine Learning Model Building
# Hyperparameters are tuned using a previous grid search result for brevity


best_rf_model = RandomForestClassifier(n_estimators=100, max_depth=None,
    min_samples_split=2, min_samples_leaf=1, random_state=42)
best_rf_model.fit(X_train, y_train)

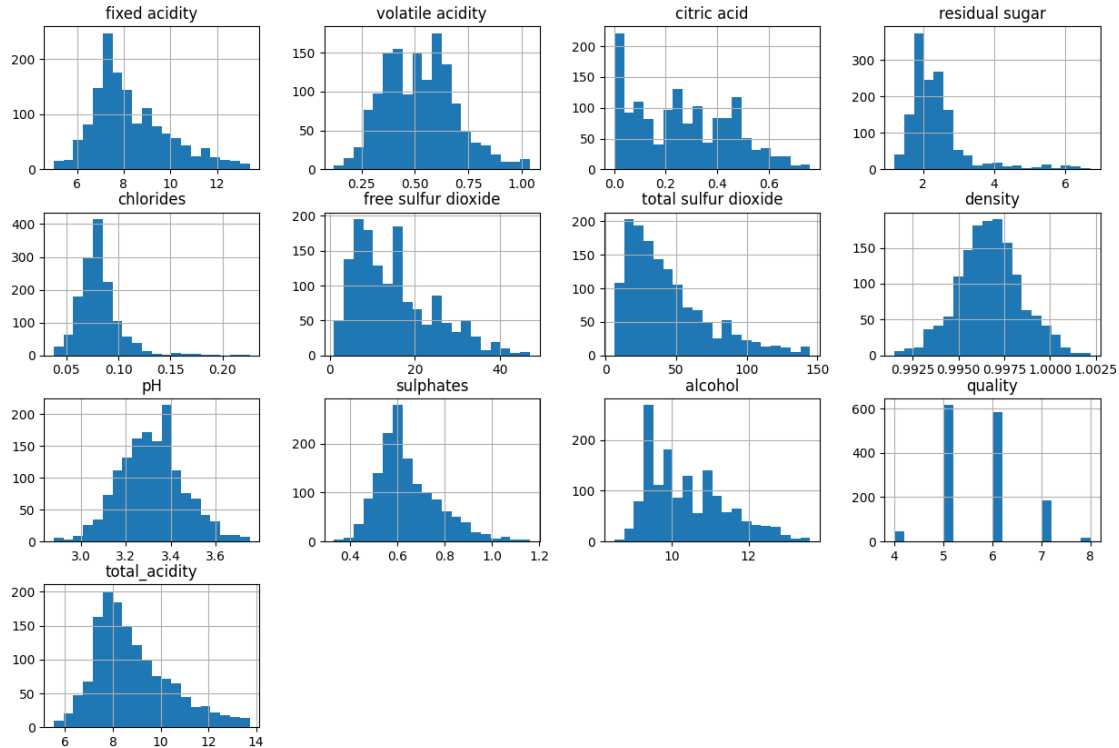
# Task 5: Evaluate the Model
y_pred = best_rf_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Display classification report for more evaluation metrics
print("Classification Report:\n", classification_report(y_test, y_pred))

# Task 6: Test with a Random Observation
random_observation = np.array([[7.0, 0.3, 0.2, 2.0, 0.075, 12.0, 30.0, 0.997, 3.
    5, 0.6, 10.0, 9.0]])
random_observation = scaler.transform(random_observation)
predicted_quality = best_rf_model.predict(random_observation)
print("Predicted wine quality for the random observation:",
    predicted_quality[0])

```





Accuracy: 0.7137931034482758

Classification Report:

	precision	recall	f1-score	support
4	0.00	0.00	0.00	10
5	0.72	0.78	0.75	118
6	0.69	0.73	0.71	123
7	0.78	0.68	0.72	37
8	0.00	0.00	0.00	2
accuracy			0.71	290
macro avg	0.44	0.44	0.44	290
weighted avg	0.69	0.71	0.70	290

Predicted wine quality for the random observation: 6

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
```

